

FastAPI RAG Project Test Document

This is a sample PDF file for testing RAG pipeline.
It contains some information about AI and Databases.

Retrieval Augmented Generation (RAG) combines:

1. Retrieval of relevant documents using a vector DB
2. Generation of answers using a Large Language Model (LLM).

Databases like PostgreSQL and MongoDB store structured data,
while Vector DBs like FAISS store embeddings for semantic search.